

# Prediction of ICU Readmissions

Unplanned readmissions to ICU contribute to high health care costs and poor patient outcomes. 6-7% of all ICU cases see a readmission within 72 hours. Machine learning models on electronic health record data can help identify these cases, providing more information about short and long-term risks to clinicians at the time of ICU discharge

## 1. Topic

In this project I'll try to explore how to predict ICU readmission based on publicly available dataset (MIMIC) and optimize the model to get a reasonable performance

## 2. Dataset


In MIMIC III dataset, (<https://mimic.physionet.org/>), we have a registry developed by the MIT Lab for Computational Physiology, comprising deidentified health data associated with ~40,000 critical care patients. It includes demographics, vital signs, laboratory tests, medications, and more.

In my github page ([https://github.com/ZepengHuo/deep\\_learning\\_class.git](https://github.com/ZepengHuo/deep_learning_class.git)), you can find the extracted data, stored in csv file format, named as df\_MASTER\_DATA.csv. It has 2354 features, ranging from heart rate, blood pressure, etc.

In the dataset, you can find out whether a patient has been readmitted in 72 hours, and I'll use that as my training labels

## 3. Neural network model

I used grid search to search all the optimal hyper-parameters (e.g. number of neurons) to train a neural network, code is ([GridSearch\\_NN.py](#)), result is ([Grid\\_search\\_results.png](#)).



As shown in the figure, I have input variables from the left and one output (readmission or not) on the right.

```
import imp
Using TensorFlow backend.
/home/grads/g/guangzhou92/enter/envs/py36/lib/python3.6/site-packages/sklearn/model_selection/_split.py:1943: FutureWarning: You should specify a value for 'cv' instead of relying on the default value. The default value will change from 3 to 5 in version 0.22.
  warnings.warn(CV_WARNING, FutureWarning)
2019-04-26 09:13:18.883789: I tensorflow/core/platform/cpu_feature_guard.cc:141] Your CPU supports instructions that this TensorFlow binary was not compiled to use: SSE4.1 SSE4.2 AVX AVX2 FMA
2019-04-26 09:13:19.361881: I tensorflow/core/common_runtime/gpu/gpu_device.cc:1411] Found device 0 with properties:
  name: GeForce GTX 1080 Ti major: 6 minor: 1 memoryClockRate(GHz): 1.582
  pciBusID: 0000:08:00:0
totalMemory: 10.92GiB freeMemory: 7.85GiB
2019-04-26 09:13:19.361848: I tensorflow/core/common_runtime/gpu/gpu_device.cc:1498] Adding visible gpu devices: 0
2019-04-26 09:13:19.819329: I tensorflow/core/common_runtime/gpu/gpu_device.cc:971] Device interconnect StreamExecutor with strength 1 edge matrix:
  0
2019-04-26 09:13:19.819377: I tensorflow/core/common_runtime/gpu/gpu_device.cc:977] 0
2019-04-26 09:13:19.819387: I tensorflow/core/common_runtime/gpu/gpu_device.cc:990] 0:  N
2019-04-26 09:13:19.819787: I tensorflow/core/common_runtime/gpu/gpu_device.cc:1183] Created TensorFlow device (/job:localhost/replica:0/task:0/device:GPU:0 with 7571 MB memory) -> physical GPU (device: 0
, name: GeForce GTX 1080 Ti, pci bus id: 0000:08:00:0, compute capability: 6.1)
Best: 0.959029 using {'neurons': 50}
0.959029 (0.001525) with: {'neurons': 50}
0.958771 (0.001726) with: {'neurons': 100}
0.959029 (0.001525) with: {'neurons': 200}
0.959029 (0.001525) with: {'neurons': 300}
0.959029 (0.001525) with: {'neurons': 400}
0.959029 (0.001525) with: {'neurons': 500}
0.959029 (0.001525) with: {'neurons': 1000}
Best: 0.959029 using {'optimizer': 'RMSprop'}
0.958888 (0.001375) with: {'optimizer': 'SGD'}
0.959029 (0.001525) with: {'optimizer': 'RMSprop'}
0.959029 (0.001525) with: {'optimizer': 'Adagrad'}
0.959029 (0.001525) with: {'optimizer': 'Adadelta'}
0.959005 (0.001543) with: {'optimizer': 'Adam'}
0.959029 (0.001525) with: {'optimizer': 'Adanamax'}
0.959029 (0.001525) with: {'optimizer': 'Nadam'}
Best: 0.959029 using {'init_mode': 'lecun_uniform'}
0.958888 (0.001628) with: {'init_mode': 'uniform'}
0.959029 (0.001525) with: {'init_mode': 'lecun_uniform'}
0.958782 (0.001557) with: {'init_mode': 'normal'}
0.959029 (0.001525) with: {'init_mode': 'zero'}
0.959005 (0.001543) with: {'init_mode': 'glorot_normal'}
0.959029 (0.001525) with: {'init_mode': 'glorot_uniform'}
0.959029 (0.001525) with: {'init_mode': 'he_normal'}
0.959029 (0.001525) with: {'init_mode': 'he_uniform'}
Best: 0.959029 using {'batch_size': 1024, 'epochs': 50}
0.959029 (0.001525) with: {'batch_size': 512, 'epochs': 50}
0.959029 (0.001525) with: {'batch_size': 512, 'epochs': 100}
0.959029 (0.001525) with: {'batch_size': 512, 'epochs': 200}
0.959029 (0.001525) with: {'batch_size': 1024, 'epochs': 50}
0.959029 (0.001525) with: {'batch_size': 1024, 'epochs': 100}
0.959029 (0.001525) with: {'batch_size': 1024, 'epochs': 200}
0.959029 (0.001525) with: {'batch_size': 2048, 'epochs': 50}
0.959029 (0.001525) with: {'batch_size': 2048, 'epochs': 100}
0.959029 (0.001525) with: {'batch_size': 2048, 'epochs': 200}
0.959029 (0.001525) with: {'batch_size': 4096, 'epochs': 50}
0.959029 (0.001525) with: {'batch_size': 4096, 'epochs': 100}
0.959029 (0.001525) with: {'batch_size': 4096, 'epochs': 200}
(py36) guangzhou92@cse-stm1-g1:~/Course/CS_Deep_learning_class/NN_HWS █
```

The grid search result is shown above. And I used the best of each hyper-parameter to train the neural network.

X\_train is 42664 rows × 2353 columns

X\_test is 10207 rows × 2353 columns

Input layer: 2353

Hidden layer 1: 50

Hidden layer 2: 100

Output layer: 1

#### 4. Annotated code

```
import os
os.environ["CUDA_DEVICE_ORDER"]="PCI_BUS_ID"
device_nb = 2
os.environ["CUDA_VISIBLE_DEVICES"]=str(device_nb)

import pandas as pd
import datetime
from xgboost import XGBClassifier
from sklearn.model_selection import GridSearchCV
import tensorflow as tf
from keras import backend as K
import tarfile
import numpy as np
import _pickle as cPickle
import os
import wfdb
from datetime import datetime
from datetime import timedelta
import numpy
from sklearn.model_selection import GridSearchCV
from keras.models import Sequential
from keras.layers import Dense
from keras.layers import Dropout
from keras.wrappers.scikit_learn import KerasClassifier
from keras.constraints import maxnorm

from sklearn.model_selection import KFold
import numpy as np
from sklearn.metrics import roc_auc_score
from sklearn.metrics import make_scorer

df_o = pd.read_csv("df_MASTER_DATA.csv")

import random
random.seed(9)
random.sample(range(1, 61089), 10000)

#sub-select 20% for testing
test_idx = random.sample(range(df_o.shape[0]), int(df_o.shape[0] * 0.2))
df_test = df_o.loc[test_idx]
train_idx = list(set([x for x in range(df_o.shape[0])] - set(test_idx)))
df_train = df_o.loc[train_idx]

best_neurons = 50
best_optimizer = 'RMSprop'
best_init_mode = 'lecun_uniform'
best_batch_size = 1024
best_epochs = 50

def create_model(neurons=best_neurons, optimizer=best_optimizer, init_mode=best_init_mode,
batch_size=best_batch_size,
```

```

        epochs=best_epochs):
# create model
model = Sequential()
model.add(Dense(neurons, input_dim=2291, activation='relu',
                kernel_initializer=init_mode, kernel_constraint=maxnorm(4)))
model.add(Dropout(0.2))
model.add(Dense(1, kernel_initializer=init_mode, activation='sigmoid'))
# Compile model
model.compile(loss='binary_crossentropy', optimizer=best_optimizer, metrics=['accuracy'])
#model.compile(loss='binary_crossentropy', optimizer='adam')
return model

#mean impute
#delete HADM_ID, SUBJECT_ID, ETHNICITY, MARITAL_STATUS, INSURANCE, RELIGION, INTIME, OUTTIME
#dummify
def preprocessing(df_labeled):

    #pre-processing, missingness > 0.5

    df_labeled_filtered = df_labeled

    for column in df_labeled_filtered.columns:
        missing_r = df_labeled[column].isnull().sum()/df_labeled.shape[0]
        if missing_r > 0.5:

            #keep arterial BP
            if not (column.startswith('Arterial_BP') or column.startswith('Mean_Arterial')):

                df_labeled_filtered = df_labeled_filtered.drop(columns=[column])

    #mean impute
    df_labeled_filtered = df_labeled_filtered.fillna(df_labeled_filtered.mean())

    #delete HADM_ID, SUBJECT_ID, ETHNICITY, MARITAL_STATUS, INSURANCE, RELIGION, INTIME, OUTTIME
    colmunns_todrop = ['Unnamed: 0', 'FIRST_CAREUNIT', 'ETHNICITY', 'MARITAL_STATUS', 'HADM_ID', 'ICUSTAY_ID',
                      'INSURANCE', 'RELIGION', 'INTIME', 'OUTTIME', 'SUBJECT_ID', 'LANGUAGE', 'Time_To_readmission',
                      'IsReadmitted_24hrs', 'IsReadmitted_48hrs', 'IsReadmitted_7days',
                      'IsReadmitted_30days', 'IsReadmitted_Bounceback']
    #don't delete ==> 'IsReadmitted_72hrs',

    for column_d in colmunns_todrop:

        try:
            df_labeled_filtered = df_labeled_filtered.drop(columns=column_d)
        except:
            pass

    #dummify
    df_labeled_filtered = pd.get_dummies(df_labeled_filtered)

    return df_labeled_filtered

```

```

def fivefold_aucroc(df_labeled_filtered):

    #Up- and down-sample for imbalance class
    X_1 = df_labeled_filtered[df_labeled_filtered['IsReadmitted_72hrs'] == 1]
    X_0 = df_labeled_filtered[df_labeled_filtered['IsReadmitted_72hrs'] == 0].iloc[:X_1.shape[0], :]

    X_all = X_0.append(X_1)

    X_all = df_labeled_filtered

    y = X_all['IsReadmitted_72hrs']
    X = X_all.drop(columns=['IsReadmitted_72hrs'])

    #X = df_labeled_filtered.drop(columns=['label'])
    #y = df_labeled_filtered['label']

    kf = KFold(n_splits=5, shuffle=True)
    kf.get_n_splits(X)

    fold_num = 1

    for train_index, test_index in kf.split(X):
        #print("TRAIN:", train_index, "TEST:", test_index)

        model = KerasClassifier(build_fn=create_model, epochs=best_epochs, batch_size=best_batch_size, verbose=0)

        model.fit(X.iloc[train_index], y.iloc[train_index])

        print('Fold %d' %fold_num, end=' ')
        fold_num += 1

        print(roc_auc_score(y.iloc[test_index], model.predict_proba(X.iloc[test_index]))[:,1]))

df_filtered = preprocessing(df_train)
fivefold_aucroc(df_filtered)

```

I have the code to run a 5 fold cross validation, and iterate through each of them, each time a 20% of data will be used as testing set, and the remaining will be used as training, until all data has been used.

## 5. Visualization/demonstration

Run the command in terminal:

**python medical\_prediction.py**

And you will see input and output window. In my project, I'm planning to use medical data registry to predict whether the patient will be readmitted to an ICU (a minor revision from last video). Since readmission to an ICU in a short time will incur unnecessary overhead and potentially influence the mortality rate. So if the model can predict if a patient will be readmitted at what time, it can largely reduce the patient risk and increase survival rate

The demo will first require you generate test patient data from a test set (data here is a mockup), and then it will predict the readmission for those patients. Lastly it will plot the result of how well the model predicts.

Github link: [https://github.com/ZepengHuo/deep\\_learning\\_class.git](https://github.com/ZepengHuo/deep_learning_class.git)

Video link: <https://youtu.be/KkjTmxNwnSs>